

EZHILAN CHINNASAMY

ezhilan03@gmail.com | +1 4693326833 | Dallas, TX, USA | Open to UK relocation
LinkedIn | GitHub | Portfolio

Data engineer combining hands-on platform delivery with cross-team leadership in payments. Designs data lakes, integrations and automation that make financial operations more reliable and analytics easier to use. Brings experience processing 2M records daily and translating business requirements into delivery across three teams.

SKILLS

- **Data Engineering:** ETL/ELT design, star-schema modelling, Airflow orchestration, data lakes, API/SFTP ingestion
- **Languages & Platforms:** SQL, Python, PostgreSQL, FastAPI, Azure Data Factory, AWS S3/EMR, PySpark
- **Analytics & Quality:** Financial reconciliation, evaluation harnesses, Grafana business-operations monitoring
- **Engineering Delivery:** Docker, Terraform, GitHub Actions CI/CD, CloudWatch (portfolio projects)

PROFESSIONAL EXPERIENCE

Pay with Spire, Business Intelligence Analyst Sep 2024 - Jun 2026 | Dallas, Texas, US

- Designed a bronze/silver/gold Azure data lake with Azure Data Factory pipelines processing 2M records daily for multi-team reporting, marketing automation and AI analytics.
- Led three-team delivery of settlement automation, translating payment-processing rules into domain architecture and coordinating vendor API/SFTP integration across operations and engineering.
- Built modular Python merchant onboarding with partner-file validation, API integration and regression tests, saving an estimated 8.5 hours of manual entry per week.
- Delivered three scheduled ADF pipelines for Zendesk, including incremental loads, moving API-based reporting to SQL and processing 500–1,200 changed tickets per daily run.
- Automated partner billing with Python and SQL, checking transaction counts, amounts and fees before generating and distributing monthly reports.
- Built a production multi-agent analytics layer over the warehouse using Google ADK and Vertex AI, enabling self-service NL2SQL/BQML querying and offloading approximately 30% of data-team workload.

The University of Texas at Dallas, Student Assistant Jun 2024 - Sep 2024

- Automated public-record integration and trained 50+ student researchers, accelerating project delivery by 30%.
- Consolidated 100k+ fragmented data points spanning five decades of Texas tech history into a timeline application adopted across departments for research.

Ernst & Young Global Delivery Services, Software Engineer Feb 2021 - Oct 2022 | Chennai, India

- Increased batch throughput by 24% with PySpark pipelines on AWS EMR processing multi-terabyte logs.
- Architected a secure migration framework from legacy multi-tenant databases to an AWS S3 data lake, strengthening IAM controls and improving query performance by 35%.

PROJECTS

Multi-Agent Financial Reconciliation Engine

- Built deterministic ledger-to-settlement matching with durable allocation guards, replay-safe ingestion, LangGraph investigation tools and human-review routing.
- Shipped an AWS batch demo processing 500 synthetic transactions and 531 settlement lines; verified backup restoration and duplicate-safe replay, with 383 transactions matched and 117 retained for review.

Fintech Hybrid RAG Engine

- Built versioned document ingestion and hybrid BM25/pgvector retrieval with validated citations and explicit abstention; verified 84 tests, seven retrieval cases and eight controlled-context grounding cases.
- Deployed an on-demand AWS demo with Docker, Terraform and GitHub Actions OIDC; verified live queries, encrypted backup restoration, dependency recovery and automatic shutdown.

EDUCATION & CERTIFICATIONS

The University of Texas at Dallas, MS in Business Analytics and Artificial Intelligence Jan 2023 - Dec 2024 | GPA: 3.7

Certifications issued by UTD: Applied Machine Learning Certificate

Kumaraguru College of Technology, Bachelors in Computer Science and Engineering Jul 2017 - Jun 2021 | GPA: 8.1

Scrum Master Accredited Certification - International Scrum Institute | Mar 2026